A 16x wider SAE shrinks the object-tracking gap (26% to 15%) but does not close it
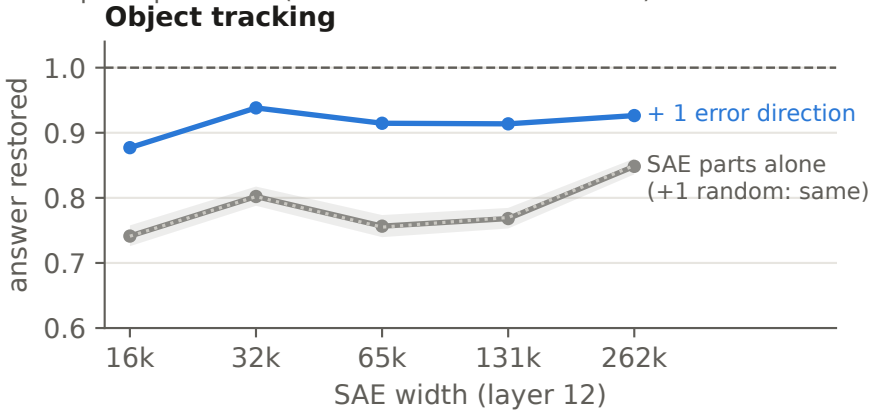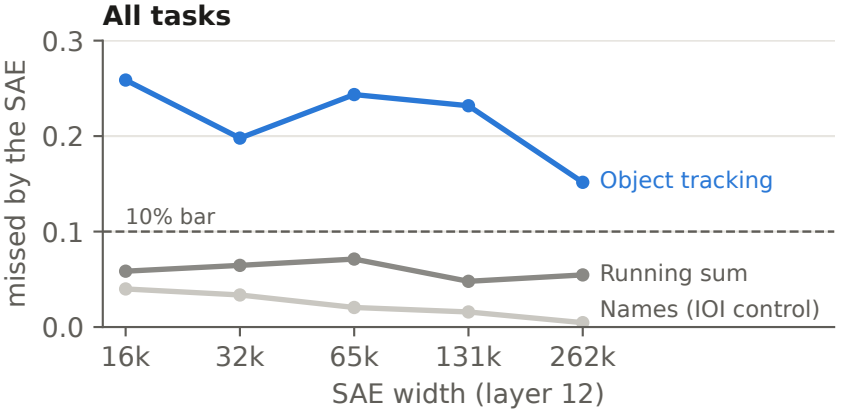
At every width one error direction recovers about half of what is missing; the other tasks' gaps are small. Note: the 131k/262k SAEs also use more active parts per token (L0 124-139 vs 75-89 at 16k)